

Bareos Python Plugins Hacking Workshop



Open Source Backup Conference

22 - 23 September 2014 | Cologne

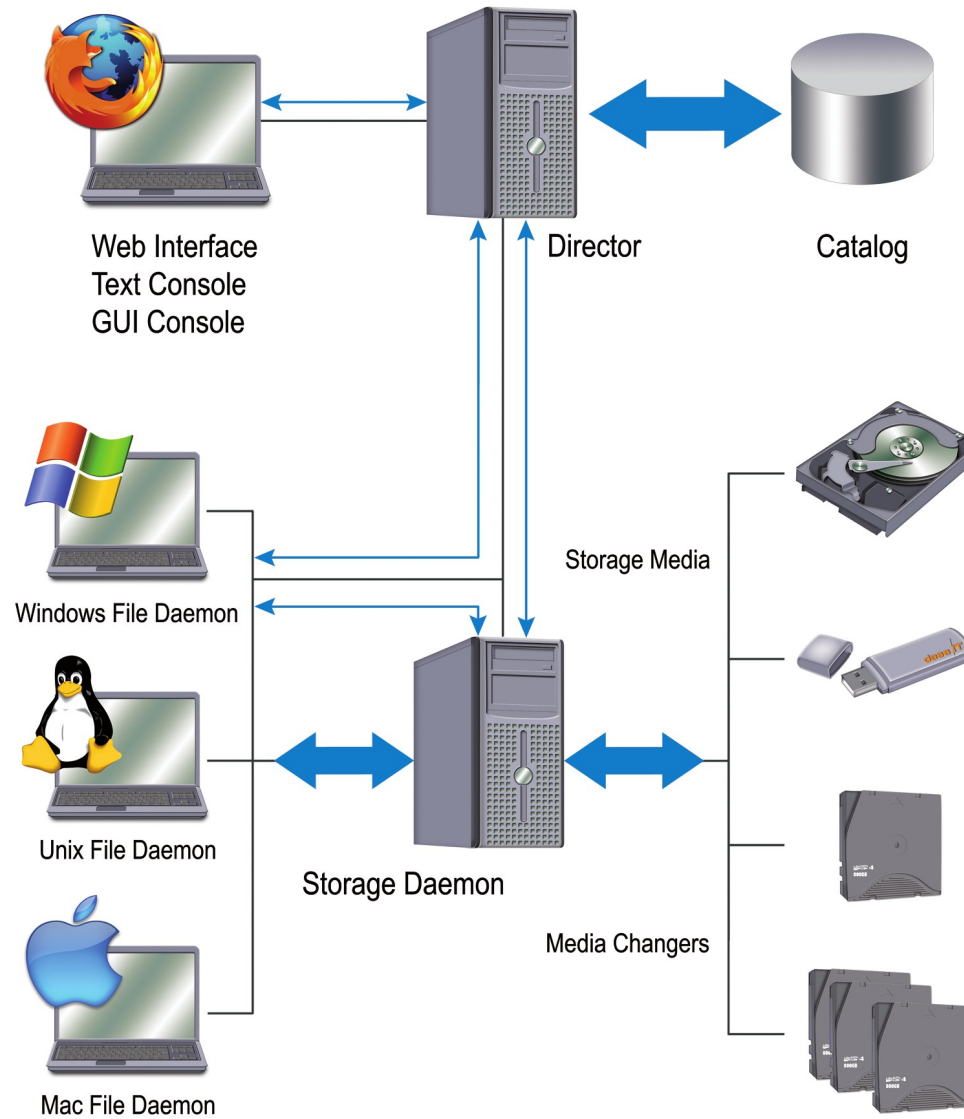
Stephan Dühr & Maik Außendorf

Sep 22, 2014

Agenda

- Bareos architecture and terminology
- Introduction
- Plugin overview (FD, SD, DIR)
- Detailed View at FileDaemon Plugins
- FD Plugin Examples
- Hacking: write your own plugin or extend existing ones
- Resume:
 - short description of work done
 - Feedback about plugin interface, questions, ideas...

Architecture Overview



Why Plugins?

- Extend Bareos functionality
 - Without touching the Bareos code
 - Can react on numerous events (in contrast to pre- and postscripts)
 - Modify Fileset
 - Extra treatment for files
 - Connect to other systems (Monitoring, Ticket, Hypervisors, Cloud, Logging, Indexer i.e. elasticsearch)
 - Application specific actions on backup and restore

C Plugin Interface

- C Coding knowledge needed
- Compiling / Linking to create a shared library
- rarely used (only well-known: bpipe)
- No known Bacula / Bareos C-Plugins for SD and Dir
- Incomplete implementation in Bacula / Bareos
- Near to zero documentation
- Legacy Python interface existed, removed from Bareos in favor of new designed interface



New Bareos Python Plugin interface

- Python knowledge wide spread among technical consultants, admins and devops
- Arbitrary Python modules available to handle a large numbers of application / APIs
- Plain Python script for FD / SD / DIR plugins
- For FD additional class based approach available
- Need Python version 2.6 or newer
- Uses distribution provided Python packages
- C code already prepared for Python 3.x



New Bareos Python Plugin interface

- Plugins configured via Bareos configuration
Pass plugin options to FD plugins
- Bareos core calls functions from the plugins on defined events
- Plugins can influence the backup process and modify Bareos variables
- Plugin usage must be explicitly enabled:
Plugin Directory = `/usr/lib/bareos/plugins`

Director Plugins

- One plugin, fixed name: `bareos-dir.py`
Hint: use symbolic link from your plugin to `bareos-dir.py`
- Configuration snippet for director resource:

```
Director {  
    ...  
    Plugin Directory = /usr/lib/bareos/plugins  
    Plugin Names = python  
}
```


Director Plugins

- Plugin registers for events it wants to be called:

```
events = [];  
events.append(bDirEventType['bDirEventJobStart']);  
events.append(bDirEventType['bDirEventJobEnd']);  
events.append(bDirEventType['bDirEventJobInit']);  
events.append(bDirEventType['bDirEventJobRun']);  
RegisterEvents(context, events);
```
- Bareos core calls function `handle_plugin_event` with event type and context
-

Director Plugins

- Context variable can be read:

```
jobId = GetValue(context, brDirVariable['bDirVarJobId']);
```

- Available variables: http://regress.bareos.org/doxygen/html/d2/d75/namespacebareos__dir__consts.html

```
bDirVarJob = 1  
bDirVarLevel = 2  
bDirVarType = 3  
bDirVarJobId = 4  
bDirVarClient = 5  
bDirVarNumVols = 6  
bDirVarPool = 7  
bDirVarStorage = 8  
bDirVarWriteStorage = 9  
bDirVarReadStorage = 10  
bDirVarCatalog = 11  
bDirVarMediaType = 12  
bDirVarJobName = 13  
bDirVarJobStatus = 14  
bDirVarPriority = 15,  
bDirVarVolumeName = 16,  
bDirVarCatalogRes = 17  
bDirVarJobErrors = 18  
bDirVarJobFiles = 19  
bDirVarSDJobFiles = 20  
bDirVarSDErrors = 21  
bDirVarFDJobStatus = 22,  
bDirVarSDJobStatus = 23,  
bDirVarPluginDir = 24  
bDirVarLastRate = 25  
bDirVarJobBytes = 26,  
bDirVarReadBytes = 27
```

Director Plugins

- Example: bareos-dir-nagios.py (stub)
 - On bDirEventJobEnd send job information to Nagios / Icinga
 - Evaluating bDirVarJobStatus

- Code excerpt:

```
def handle_plugin_event(context, event):
    ...
    elif event == bDirEventType['bDirEventJobEnd']:
        jobName = GetValue(context, brDirVariable['bDirVarJobName']);
        jobStatus = chr(GetValue(context, brDirVariable['bDirVarJobStatus']))
        ...
        if (jobStatus == 'E' or jobStatus == 'f'):
            nagiosResult = 2; # critical
            nagiosMessage = "CRITICAL: Bareos job %s on %s with id %d failed with %s errors (%s), %d jobBytes" %
(jobName,jobClient,jobId,jobErrors,jobStatus,jobBytes);
# send nagiosResult and nagiosMessage to Icinga / Nagios
```

- Send to Nagios / Icinga needs to be adapted
https://github.com/bareos/bareos-contrib/tree/master/dir-plugins/nagios_icinga

SD Plugins

- One plugin, fixed name: `bareos-sd.py`
Hint: use symbolic link from your plugin to `bareos-sd.py`
- Configuration snippet to enable plugins in Storage resource of `bareos-sd.conf`:
Storage {
 ...
 Plugin Directory = /usr/lib64/bareos/plugins
 Plugin Names = python
}

SD Plugins

- when omitting `Plugin Names`, `sd` will load all plugins in plugin directory matching glob `*-sd.so`
- explicitly specifying multiple plugins:
`Plugin Names = python:autoxflate`
will load `python-sd.so` and `autoxflate-sd.so`
- registering events and variable handling like `dir` plugin
- until now, only variables `bsdVarJobId`, `bsdVarJobName` and `bsdVarPluginDir` are exposed
(`python-sd.c` still in early development state)
- future: all variables as described should be available, see http://regress.bareos.org/doxygen/html/d2/de6/namespacebareos__sd__consts.html

- example: bareos-sd.py

```
from bareosd import *
from bareos_sd_consts import *

def load_bareos_plugin(context):
    DebugMessage(context, 100, "load_bareos_plugin called\n")
    events = []
    events.append(bsdEventType['bsdEventJobStart'])
    events.append(bsdEventType['bsdEventJobEnd'])
    RegisterEvents(context, events)
    return bRCs['bRC_OK']

def handle_plugin_event(context, event):
    if event == bsdEventType['bsdEventJobStart']:
        DebugMessage(context, 100, "bsdEventJobStart event triggered\n")
        jobname = GetValue(context, bsdrVariable['bsdVarJobName'])
        DebugMessage(context, 100, "Job " + jobname + " starting\n")
    ...
```

FD Plugins

- how to enable Python Plugins in FD?
- install `bareos-filedaemon-python-plugin`
- in `bareos-fd.conf` add or uncomment:

```
FileDaemon {  
    ...  
    Plugin Directory = /usr/lib64/bareos/plugins  
    Plugin Names = python  
    ...  
}
```
- like for SD and Dir Plugins, `Plugin Names` can be omitted. Then all Plugins matching glob `*-fd.so` will be loaded

FD Plugins

- multiple plugins possible
- the `Plugin` parameter in Director's FileSet resource determines which python plugin is used with which parameters.

Syntax:

```
Plugin = python:module_path=<path-to-python-modules>;module_name=<python-module-to-load>;<custom-param1>=<custom-value1>;...
```

- `module_path` and `module_name` are mandatory (used by `python-fd.so`)
- anything else is arbitrary, the complete string is passed to the hook function `parse_plugin_definitions()`
- two Plugin-Types:
Command-Plugins and Option-Plugins

How do FD Plugins work (1)

- When a Job is run, Director passes plugin definition to FD, eg. `module_path=/usr/lib64/bareos/plugins:module_name=bareos-fd`

FD (python-fd.so) does the following:

- instantiates new Python interpreter
- extends the Python search path with the given `module_path`
- imports the module given by `module_name` (for the example above, would be `bareos-fd.py`)
- makes callback methods available for Python, use `from bareosfd import *` in Python code

How do FD Plugins work (2)

- calls `load_bareos_plugin()` in the python plugin code
- calls `parse_plugin_definition(context, plugindef)` in the python code
 - `plugindef` is the complete string as configured in Director (Plugin = ...), to be parsed by python code
- different processing loop depending on type of Plugin (Command/Option)



FD Command-Plugin Configuration

- Command Plugin Configuration in Include section of FileSet Resource in bareos-dir.conf:

```
FileSet {  
    Name = "test_PyLocalFileset_Set"  
    Include {  
        Plugin =  
        "python:module_path=/usr/lib64/bareos/plugins:module_name=bareos-fd-local-fileset:filename=/tmp/datafile"  
    }  
}
```

FD Option-Plugin Configuration

- Option Plugin Configuration in Options section of Include Section of FileSet Resource in bareos-dir.conf:

```
FileSet {  
    Name = "test_PyOptionInteract_Set"  
    Include {  
        File = /data/project_1  
        Options {  
            Plugin =  
"python:module_path=/usr/lib64/bareos/plugins:module_name=bareos-  
fd-file-interact"  
        }  
    }  
}
```

- Note: for Option-Plugin must define files to backup using File = ... while for Command-Plugin need not



Difference FD Command-/Option-Plugins (1)

- Major Difference:
 - Command-Plugin determines what is being backed up, must also handle Diff/Inc itself
 - Option-Plugin gets which files to backup based on whats configured in Director, Diff/Inc handling done by FD



Difference FD Command-/Option-Plugins (2)

- Command-Plugin processing
 - `start_backup_file(context, savepkt)` must set `savepkt` properties for each file to back up
 - `plugin_io(context, IOP)` must handle IO Operations
 - Backup: `open(r)`, `read`, `close`
 - `end_backup_file(context)`
 - must return `bRCs['bRC_More']` if more files to backup
 - must return `bRCs['bRC_OK']` to finish the looping
 - `handle_backup_file()` is not called



Difference FD Command-/Option-Plugins (3)

- Option-Plugin processing
 - `handle_backup_file(context, savepkt)` called for each file to be processed, `savepkt` defined by FD
 - `plugin_io()` handling in the same manner as for Command-Plugin
 - `start_backup_file()` and `end_backup_file()` are not called



FD Plugins – Callback Functions

- Functions provided by python-fd.so that can be called from Python code, enabled by
`from bareosfd import *`
- Complete list: see http://regress.bareos.org/doxygen/html/d5/d0e/python-fd_8h_source.html
- Most important callback functions:
 - `JobMessage()`: Error-/Info-/Warning-Messages
 - are passed to Director, appear in messages and logs
 - `DebugMessage()`: Debug-Messages with numeric level
 - only visible when running FD in debug-mode with `-d <level>`
 - `GetValue()`: used to get variables from FD



FD Plugins – Class Based Approach

- Python FD Plugin can be monolithic
- Better: use classes and inheritance to reuse existing code easier and reduce code redundancy
- To support this approach, the package `bareos-filedaemon-python-plugin` package provides
 - `BareosFdPluginBaseclass.py`
 - Parent Class to inherit from
 - `BareosFdWrapper.py`
 - defines all functions a plugin needs and “wraps” them to the corresponding methods in the plugin class
 - a Plugin entry-point module glues them together

Messaging

- DebugMessage: Debug only
 - `DebugMessage(context, level, "message\n");`
 - context: used to pass information from core to plugin, don't touch
 - Level: Debug Level, use ≥ 100
 - Sample:
`DebugMessage(context, 100, "handle_backup_file called with " + str(savepkt) + "\n");`

Messaging

- JobMessage: Sent to messaging system
 - `JobMessage(context, bJobMessageType, "Message\n");`
 - Type: Controls job result, M_INFO, M_ERROR, M_WARNING, M_ABORT
http://regress.bareos.org/doxygen/html/dd/dbb/namespacebareos__fd__consts.html
 - Sample:
`JobMessage(context, bJobMessageType['M_INFO'], "Option Plugin file interact on" + savepkt.fname + "\n");`

Return Codes

- Return Codes control processing, no impact on overall job status.
- Depending on context / function

- Use consts:

```
return bRCs['bRC_OK'];  
return bRCs['bRC_Skip']; # skips current file  
return bRCs['bRC_Error']; # error but continue  
return bRCs['bRC_More']; # in end_backup_file, more files to  
backup  
...
```



FD Plugin: bareos-fd-local-fileset.py

- Reads a local file on fd with filenames to backup
 - Demonstration / template plugin, functionality can be achieved better by fileset configuration:

File = "\\</localfile/on/client"

- Configuration in fileset resource as command plugin (extends fileset):

```
Plugin = "python:module_path=/usr/lib64/bareos/plugins:module_name=bareos-fd-local-fileset:filename=/tmp/datafile"
```

- Plugin: /usr/lib64/bareos/plugins/bareos-fd-local-fileset.py

Code excerpt:

```
from bareosfd import *
from bareos_fd_consts import *
from BareosFdWrapper import *
from BareosFdPluginLocalFileset import *
def load_bareos_plugin(context, plugindef):
    BareosFdWrapper.bareos_fd_plugin_object = BareosFdPluginLocalFileset (context, plugindef);
    return bRCs['bRC_OK'];
```

- Rest is done in class BareosFdPluginLocalFileset

- Class inherits from BareosFdPluginBaseclass
- Method `parse_plugin_definition`
Parses the options, filename is mandatory
Reads filenames from file into array
`self.files_to_backup`
- Method `start_backup_file` asks plugin, if there is anything to backup, sets `savepkt`:

```
file_to_backup = self.files_to_backup.pop();  
savepkt.fname = file_to_backup;  
savepkt.type = bFileType['FT_REG'];  
return bRCs['bRC_OK'];
```

- Method `end_backup_file` called to ask plugin if there is more to backup:

```
if self.files_to_backup:
    # there is more to backup, go to start_backup_file again
    return bRCs['bRC_More'];
else
    # no more to backup from this plugin, done
    return bRCs['bRC_OK'];
```

- Basic IO operations covered in base class
 - Method `plugin_io` handles file read / write operations

- For restore: some more things to do

- Directories have to be created

```
def create_file (self, context, restorepkt):  
    FNAME = restorepkt.ofname;  
    dirname = os.path.dirname (FNAME);  
    if not os.path.exists(dirname):  
        os.makedirs(dirname);  
    open (FNAME, 'wb').close();  
    restorepkt.create_status = bCFs['CF_EXTRACT'];  
    return bRCs['bRC_OK'];
```

- Similar in method `plugin_io` for writing

- Overload this method in your class, if you need different handling



FD Plugin: bareos-fd-file-interact.py

- Stub / example for option plugin
- Gets called before a file is read for backup, giving the possibility to do anything with this file
- Configuration in fileset resource as option plugin in options section:

```
Plugin = "python:module_path=/usr/lib64/bareos/plugins:module_name=bareos-fd-file-interact"
```

- Plugin: /usr/lib64/bareos/plugins/bareos-fd-local-fileset.py

Uses BareosFdWrapper with class:

BareosFdPluginFileInteract



BareosFdPluginFileInteract

- Method `handle_backup_file` called for each file to backup with param `savepkt`
 - Backup file with full path in `savepkt.fname`
 - File type in `savepkt.type`

- File types (selection):

`FT_REG = 3 # Regular file`

`FT_LNK = 4 # Soft link`

`FT_FIFO = 17 # Raw FiFo device`

Full list: http://regress.bareos.org/doxygen/html/d4/d6f/filetypes_8h.html

- `FT_RESTORE_FIRST = 25`

Special Restore object, can contain metadata needed for restore. Needs to be created as virtual file, these objects will be restored first, handled by your plugin.

My SQL Plugin

- FD Plugin for MySQL Backup contributed by Evan Felix (<https://github.com/karcaw>)
- Available at <https://github.com/bareos/bareos-contrib/tree/master/fd-plugins/mysql-python>
- runs `mysql -B -N -e 'show databases'` to get the list of databases to back up
- runs `mysqldump %s --events --single-transaction | gzip` for each database, using `os.popen()` (pipe)
- `plugin_io()` reads the pipe, no temporary local disk space needed for the dump

My SQL Plugin

- Possible enhancements:
 - add option to specify which DB(s) to backup
 - add option to specify compression program
 - pbzip2 or pigz may accelerate the backup, if system has multiple CPUs/Cores
 - add the ability to disable compression by pipe, to have Bareos handle compression

VMware plugin

- Work in progress to allow Snapshot based VM Backups, not yet published
- Use VMware's CBT (Changed Block Tracking) to allow space efficient Full and Incremental Backups
- Coping with the complex VMware API is not easy
- Using Java was a no-go for Bareos
- Until December 2013 several more or less useful Projects to use the API with Python were around: PySphere, Psphere, PyVISDK
- in December 2013 pyvmomi appeared on github, a Python SDK for the Vsphere API, sponsored/supported by VMware

VMware plugin

- Also requires using the Virtual Disk Development Kit (VDDK), it's a collection of C Libraries, sometimes also named vmware-vix-disklib
- No good or VMware sponsored/supported VDDK binding for Python exists
- Using <https://github.com/xuru/vixDiskLib> in a Python FD Plugin failed because VDDK comes with some older libs that caused unresolvable conflicts/errors
- New approach uses a separate program developed in C that handles VDDK and pipes data to the FD Plugin

VMware plugin

- What works already:
 - Creating/Removing a snapshot
 - CBT base dump (Full), CBT advantage: only active portions of virtual disks dumped
 - manual restore
- Still Missing:
 - Saving VM Metadata/Config, needed to re-create the same VM or a identically configured VM for restore
 - Creating a VM for restore and the restore process itself
 - Citation from http://pubs.vmware.com/vsphere-55/index.jsp#com.vmware.vddk.pg.doc/vddkBkupVadp.9.4.html?path=7_4_0_6_2_0#1020552
“This section shows how to create a VirtualMachine object, which is complicated but necessary so you can restore data into it.”
 - Differential/Incremental Backup/Restore

- Group together (2/3 people per group)
- Get one of the existing plugins up and running
- Extend existing plugin, e.g.
 - Icinga / Director plugin: configurable interface to Nagios / Icinga (send_nsca...)
 - Mysql: make databases to backup configurable, gzip optional, restore directly to db optional
 - Local Fileset plugin: directories, optionally include / exclude files belonging to a specific user

- More ideas:
 - Director plugin: connect to ticket system (otrs, rt)
 - FD option plugin: pass files to elasticsearch for indexing plus backup meta information
 - FD option plugin: create local log for every file in backup with timestamp and checksum
 - FD option plugin: gpg encrypt every file on the fly

Live Hacking

- More ideas – application specific plugins
 - IMAP / Cyrus: restore to specific mailbox directory
 - Open Xchange (backup / restore of single objects)
 - Kolab
 - other SQL or NoSQL Databases
 - oVirt/RHEV Backup (REST API now has backup functions)
 - Snapshot based KVM (some ideas next slide)
 - Other applications?

- Ideas regarding KVM Backup
 - KVM/qemu has nothing like VMware CBT
 - Promising design proposals like <http://wiki.qemu.org/Features/Livebackup> have never been completed
 - a CBT-like approach using external QCOW2 snapshots/overlays could be derived from <https://kashyapc.fedorapeople.org/virt/lc-2012/snapshots-handout.html>
 - Guest-Agent quiescing actions should be looked at
 - Performance impact of overlay chaining?
 - use python libvirt bindings



Contact and links

- Subscription, Support, References, Partner:
<http://www.bareos.com>
- Community, Documentation, Download:
<http://www.bareos.org>
- GIT Bareos:
<https://github.com/bareos>
- GIT Bareos contrib for plugins:
<https://github.com/bareos/bareos-contrib>
- Bug- and feature- tracker Mantis:
<https://bugs.bareos.org>